

CSE544: SQL

Monday 3/27 and Wednesday 3/29, 2006

SQL Introduction

Standard language for querying and manipulating data

Structured Query Language

Many standards out there:

- ANSI SQL, SQL92 (a.k.a. SQL2), SQL99 (a.k.a. SQL3),
- Vendors support various subsets: watch for fun discussions in class !

SQL

- Data Definition Language (DDL)
 - Create/alter/delete tables and their attributes
 - Following lectures...
- Data Manipulation Language (DML)
 - Query one or more tables – discussed next !
 - Insert/delete/modify tuples in tables

Table name

Attribute names

Tables in SQL

Product

PName	Price	Category	Manufacturer
Gizmo	\$19.99	Gadgets	GizmoWorks
Powergizmo	\$29.99	Gadgets	GizmoWorks
SingleTouch	\$149.99	Photography	Canon
MultiTouch	\$203.99	Household	Hitachi

Tuples or rows

Tables Explained

- The *schema* of a table is the table name and its attributes:

Product(PName, Price, Category, Manufacturer)

- A *key* is an attribute whose values are unique; we underline a key

Product(PName, Price, Category, Manufacturer)

Data Types in SQL

- Atomic types:
 - Characters: CHAR(20), VARCHAR(50)
 - Numbers: INT, BIGINT, SMALLINT, FLOAT
 - Others: MONEY, DATETIME, ...
- Every attribute must have an atomic type
 - Hence tables are flat
 - Why ?

Tables Explained

- A tuple = a record
 - Restriction: all attributes are of atomic type
- A table = a set of tuples
 - Like a list...
 - ...but it is unordered:
no **first()**, no **next()**, no **last()**.

SQL Query

Basic form: (plus many many more bells and whistles)


```
SELECT <attributes>  
FROM   <one or more relations>  
WHERE  <conditions>
```


Simple SQL Query

Product

PName	Price	Category	Manufacturer
Gizmo	\$19.99	Gadgets	GizmoWorks
Powergizmo	\$29.99	Gadgets	GizmoWorks
SingleTouch	\$149.99	Photography	Canon
MultiTouch	\$203.99	Household	Hitachi

```
SELECT *  
FROM Product  
WHERE category='Gadgets'
```



“selection”


PName	Price	Category	Manufacturer
Gizmo	\$19.99	Gadgets	GizmoWorks
Powergizmo	\$29.99	Gadgets	GizmoWorks

Simple SQL Query

Product

PName	Price	Category	Manufacturer
Gizmo	\$19.99	Gadgets	GizmoWorks
Powergizmo	\$29.99	Gadgets	GizmoWorks
SingleTouch	\$149.99	Photography	Canon
MultiTouch	\$203.99	Household	Hitachi

```
SELECT PName, Price, Manufacturer
FROM   Product
WHERE  Price > 100
```



“selection” and
“projection”

PName	Price	Manufacturer
SingleTouch	\$149.99	Canon
MultiTouch	\$203.99	Hitachi


Notation



Input Schema

Product(PName, Price, Category, Manufacturer)

```
SELECT PName, Price, Manufacturer
FROM   Product
WHERE  Price > 100
```



Answer(PName, Price, Manufacturer)



Output Schema

Details

- Case insensitive:
 - Same: SELECT Select select
 - Same: Product product
 - Different: 'Seattle' 'seattle'
- Constants:
 - 'abc' - yes
 - "abc" - no


The **LIKE** operator

```
SELECT *  
FROM Products  
WHERE PName LIKE '%gizmo%'
```

- s **LIKE** p: pattern matching on strings
- p may contain two special symbols:
 - % = any sequence of characters
 - _ = any single character

Eliminating Duplicates


```
SELECT DISTINCT category  
FROM Product
```



Category
Gadgets
Photography
Household

Compare to:

```
SELECT category  
FROM Product
```



Category
Gadgets
Gadgets
Photography
Household

Ordering the Results


```
SELECT pname, price, manufacturer  
FROM Product  
WHERE category='gizmo' AND price > 50  
ORDER BY price, pname
```

Ties are broken by the second attribute on the ORDER BY list, etc.

Ordering is ascending, unless you specify the DESC keyword.


PName	Price	Category	Manufacturer
Gizmo	\$19.99	Gadgets	GizmoWorks
Powergizmo	\$29.99	Gadgets	GizmoWorks
SingleTouch	\$149.99	Photography	Canon
MultiTouch	\$203.99	Household	Hitachi

```
SELECT DISTINCT category
FROM Product
ORDER BY category
```




?

```
SELECT Category
FROM Product
ORDER BY PName
```



?

```
SELECT DISTINCT category
FROM Product
ORDER BY PName
```



?

Keys and Foreign Keys

Company



<u>CName</u>	StockPrice	Country
GizmoWorks	25	USA
Canon	65	Japan
Hitachi	15	Japan

Product

<u>PName</u>	Price	Category	Manufacturer
Gizmo	\$19.99	Gadgets	GizmoWorks
Powergizmo	\$29.99	Gadgets	GizmoWorks
SingleTouch	\$149.99	Photography	Canon
MultiTouch	\$203.99	Household	Hitachi



Foreign
key

Joins

Product (pname, price, category, manufacturer)

Company (cname, stockPrice, country)

Find all products under \$200 manufactured in Japan;
return their names and prices.

```
SELECT PName, Price  
FROM Product, Company  
WHERE Manufacturer=CName AND Country='Japan'  
AND Price <= 200
```



Join
between Product
and Company

Joins


Product

PName	Price	Category	Manufacturer
Gizmo	\$19.99	Gadgets	GizmoWorks
Powergizmo	\$29.99	Gadgets	GizmoWorks
SingleTouch	\$149.99	Photography	Canon
MultiTouch	\$203.99	Household	Hitachi

Company

Cname	StockPrice	Country
GizmoWorks	25	USA
Canon	65	Japan
Hitachi	15	Japan

```
SELECT PName, Price
FROM Product, Company
WHERE Manufacturer=CName AND Country='Japan'
AND Price <= 200
```



PName	Price
SingleTouch	\$149.99

More Joins

Product (pname, price, category, manufacturer)

Company (cname, stockPrice, country)

Find all Chinese companies that manufacture products both in the 'electronic' and 'toy' categories

```
SELECT  cname
```

```
FROM
```

```
WHERE
```

A Subtlety about Joins

Product (pname, price, category, manufacturer)

Company (cname, stockPrice, country)

Find all countries that manufacture some product in the 'Gadgets' category.

```
SELECT Country
FROM Product, Company
WHERE Manufacturer=CName AND Category='Gadgets'
```

Unexpected duplicates

A Subtlety about Joins


Product

<u>Name</u>	Price	Category	Manufacturer
Gizmo	\$19.99	Gadgets	GizmoWorks
Powergizmo	\$29.99	Gadgets	GizmoWorks
SingleTouch	\$149.99	Photography	Canon
MultiTouch	\$203.99	Household	Hitachi

Company

<u>Cname</u>	StockPrice	Country
GizmoWorks	25	USA
Canon	65	Japan
Hitachi	15	Japan

```
SELECT Country
FROM Product, Company
WHERE Manufacturer=CName AND Category='Gadgets'
```



Country
??
??

What is
the problem ?
What's the
solution ?

Tuple Variables

Person(pname, address, worksfor)

Company(cname, address)

```
SELECT DISTINCT pname, address
FROM      Person, Company
WHERE     worksfor = cname
```

Which
address ?



```
SELECT DISTINCT Person.pname, Company.address
FROM      Person, Company
WHERE     Person.worksfor = Company.cname
```



```
SELECT DISTINCT x.pname, y.address
FROM      Person AS x, Company AS y
WHERE     x.worksfor = y.cname
```

Meaning (Semantics) of SQL Queries

```
SELECT a1, a2, ..., ak  
FROM R1 AS x1, R2 AS x2, ..., Rn AS xn  
WHERE Conditions
```

```
Answer = {}  
for x1 in R1 do  
  for x2 in R2 do  
    .....  
    for xn in Rn do  
      if Conditions  
        then Answer = Answer ∪ {(a1, ..., ak)}  
return Answer
```


An Unintuitive Query

```
SELECT DISTINCT R.A  
FROM R, S, T  
WHERE R.A=S.A OR R.A=T.A
```

What does it compute ?

Computes $R \cap (S \cup T)$

But what if $S = \phi$?

Subqueries Returning Relations

Company(name, city)

Product(pname, maker)

Purchase(id, product, buyer)

Return cities where one can find companies that manufacture products bought by Joe Blow

```
SELECT Company.city
FROM Company
WHERE Company.name IN
    (SELECT Product.maker
     FROM Purchase , Product
     WHERE Product.pname=Purchase.product
     AND Purchase .buyer = 'Joe Blow');
```

Subqueries Returning Relations

Is it equivalent to this ?

```
SELECT Company.city  
FROM    Company, Product, Purchase  
WHERE   Company.name= Product.maker  
        AND Product.pname = Purchase.product  
        AND Purchase.buyer = 'Joe Blow'
```

Beware of duplicates !

Removing Duplicates

```
SELECT DISTINCT Company.city
FROM    Company
WHERE   Company.name IN
        (SELECT Product.maker
         FROM   Purchase , Product
         WHERE  Product.pname=Purchase.product
              AND Purchase .buyer = 'Joe Blow');
```

```
SELECT DISTINCT Company.city
FROM    Company, Product, Purchase
WHERE   Company.name= Product.maker
        AND Product.pname = Purchase.product
        AND Purchase.buyer = 'Joe Blow'
```

Now
they are
equivalent

Subqueries Returning Relations

You can also use: $s > \text{ALL } R$
 $s > \text{ANY } R$
 $\text{EXISTS } R$

Product (pname, price, category, maker)

Find products that are more expensive than all those produced
By “Gizmo-Works”

```
SELECT name
FROM   Product
WHERE  price > ALL (SELECT price
                    FROM   Purchase
                    WHERE  maker='Gizmo-Works')
```

Question for Database Fans and their Friends

- Can we express this query as a single `SELECT-FROM-WHERE` query, without subqueries ?

Question for Database Fans and their Friends

- Answer: all SFW queries are **monotone** (figure out what this means).
A query with **ALL** is not monotone

Correlated Queries

Movie (title, year, director, length)

Find movies whose title appears more than once.

```
SELECT DISTINCT title
FROM   Movie AS x
WHERE  year <> ANY
      (SELECT year
       FROM   Movie
       WHERE  title = x.title);
```

correlation



Note (1) scope of variables (2) this can still be expressed as single SFW

Complex Correlated Query

Product (pname, price, category, maker, year)

- Find products (and their manufacturers) that are more expensive than all products made by the same manufacturer before 1972

```
SELECT DISTINCT pname, maker
FROM   Product AS x
WHERE  price > ALL (SELECT price
                    FROM   Product AS y
                    WHERE  x.maker = y.maker AND y.year < 1972);
```

Very powerful ! Also much harder to optimize.

Aggregation

```
SELECT avg(price)
FROM Product
WHERE maker="Toyota"
```

```
SELECT count(*)
FROM Product
WHERE year > 1995
```

SQL supports several aggregation operations:

sum, count, min, max, avg

Except count, all aggregations apply to a single attribute

Aggregation: Count

COUNT applies to duplicates, unless otherwise stated:

```
SELECT Count(category)
FROM Product
WHERE year > 1995
```

same as Count(*)

We probably want:

```
SELECT Count(DISTINCT category)
FROM Product
WHERE year > 1995
```

More Examples

Purchase(product, date, price, quantity)

```
SELECT Sum(price * quantity)
FROM   Purchase
```


```
SELECT Sum(price * quantity)
FROM   Purchase
WHERE  product = 'bagel'
```

What do
they mean ?

Purchase Simple Aggregations

Product	Date	Price	Quantity
Bagel	10/21	1	20
Banana	10/3	0.5	10
Banana	10/10	1	10
Bagel	10/25	1.50	20

```
SELECT Sum(price * quantity)
FROM Purchase
WHERE product = 'bagel'
```



50 (= 20+30)

Grouping and Aggregation

Purchase(product, date, price, quantity)

Find total sales after 10/1/2005 per product.

```
SELECT    product, Sum(price*quantity) AS TotalSales
FROM      Purchase
WHERE     date > '10/1/2005'
GROUP BY  product
```

Let's see what this means...

Grouping and Aggregation


1. Compute the **FROM** and **WHERE** clauses.
2. Group by the attributes in the **GROUPBY**
3. Compute the **SELECT** clause: grouped attributes and aggregates.

1&2. FROM-WHERE-GROUPBY

Product	Date	Price	Quantity
Bagel	10/21	1	20
Bagel	10/25	1.50	20
Banana	10/3	0.5	10
Banana	10/10	1	10

3. SELECT

Product	Date	Price	Quantity
Bagel	10/21	1	20
Bagel	10/25	1.50	20
Banana	10/3	0.5	10
Banana	10/10	1	10



Product	TotalSales
Bagel	50
Banana	15

```
SELECT    product, Sum(price*quantity) AS TotalSales
FROM      Purchase
WHERE     date > '10/1/2005'
GROUP BY product
```

GROUP BY v.s. Nested Quereis

```
SELECT    product, Sum(price*quantity) AS TotalSales
FROM      Purchase
WHERE     date > '10/1/2005'
GROUP BY  product
```

```
SELECT DISTINCT x.product, (SELECT Sum(y.price*y.quantity)
                              FROM    Purchase y
                              WHERE    x.product = y.product
                              AND y.date > '10/1/2005')
          AS TotalSales
FROM      Purchase x
WHERE     x.date > '10/1/2005'
```

Another Example

What does
it mean ?

```
SELECT    product,  
          sum(price * quantity) AS SumSales  
          max(quantity) AS MaxQuantity  
FROM      Purchase  
GROUP BY product
```

HAVING Clause

Same query, except that we consider only products that had at least 100 buyers.

```
SELECT    product, Sum(price * quantity)
FROM      Purchase
WHERE     date > '10/1/2005'
GROUP BY  product
HAVING    Sum(quantity) > 30
```

HAVING clause contains conditions on aggregates.

General form of Grouping and Aggregation

SELECT S
FROM R_1, \dots, R_n
WHERE C1
GROUP BY a_1, \dots, a_k
HAVING C2



S = may contain attributes a_1, \dots, a_k and/or any aggregates but NO OTHER ATTRIBUTES

C1 = is any condition on the attributes in R_1, \dots, R_n

C2 = is any condition on aggregate expressions

General form of Grouping and Aggregation

```
SELECT  S  
FROM    R1,...,Rn  
WHERE   C1  
GROUP BY a1,...,ak  
HAVING  C2
```

Evaluation steps:

1. Evaluate FROM-WHERE, apply condition C1
2. Group by the attributes a_1, \dots, a_k
3. Apply condition C2 to each group (may have aggregates)
4. Compute aggregates in S and return the result

Advanced SQLizing


1. Getting around INTERSECT and EXCEPT
2. Quantifiers
3. Aggregation v.s. subqueries

INTERSECT and EXCEPT: not in SQL Server

1. INTERSECT and EXCEPT:


If R, S have no duplicates, then can write without subqueries (HOW ?)

```
(SELECT R.A, R.B  
FROM R)  
INTERSECT  
(SELECT S.A, S.B  
FROM S)
```



```
SELECT R.A, R.B  
FROM R  
WHERE  
EXISTS(SELECT *  
FROM S  
WHERE R.A=S.A and R.B=S.B)
```

```
(SELECT R.A, R.B  
FROM R)  
EXCEPT  
(SELECT S.A, S.B  
FROM S)
```



```
SELECT R.A, R.B  
FROM R  
WHERE  
NOT EXISTS(SELECT *  
FROM S  
WHERE R.A=S.A and R.B=S.B)
```


2. Quantifiers

Product (pname, price, company)

Company(cname, city)

Find all companies that make some products with price < 100

```
SELECT DISTINCT Company.cname
FROM    Company, Product
WHERE   Company.cname = Product.company and Product.price < 100
```

Existential: easy ! 😊

2. Quantifiers

Product (pname, price, company)

Company(cname, city)

Find all companies that make only products with price < 100

same as:

Find all companies s.t. all of their products have price < 100

Universal: hard ! ☹️

2. Quantifiers

1. Find *the other* companies: i.e. s.t. some product ≥ 100

```
SELECT DISTINCT Company.cname
FROM   Company
WHERE  Company.cname IN (SELECT Product.company
                        FROM Product
                        WHERE Produc.price  $\geq$  100)
```

2. Find all companies s.t. all their products have price < 100

```
SELECT DISTINCT Company.cname
FROM   Company
WHERE  Company.cname NOT IN (SELECT Product.company
                        FROM Product
                        WHERE Produc.price  $\geq$  100)
```

3. Group-by v.s. Nested Query

Author(login,name)

Wrote(login,url)

- Find authors who wrote ≥ 10 documents:
- Attempt 1: with nested queries

This is
SQL by
a novice

```
SELECT DISTINCT Author.name
FROM      Author
WHERE     count(SELECT Wrote.url
                  FROM Wrote
                  WHERE Author.login=Wrote.login)
          > 10
```

3. Group-by v.s. Nested Query

- Find all authors who wrote at least 10 documents:
- Attempt 2: SQL style (with GROUP BY)

```
SELECT    Author.name  
FROM      Author, Wrote  
WHERE     Author.login=Wrote.login  
GROUP BY  Author.name  
HAVING    count(wrote.url) > 10
```



This is
SQL by
an expert

No need for **DISTINCT**: automatically from **GROUP BY**

3. Group-by v.s. Nested Query

Author(login,name)

Wrote(login,url)

Mentions(url,word)

Find authors with vocabulary ≥ 10000
words.

```
SELECT    Author.name
FROM      Author, Wrote, Mentions
WHERE     Author.login=Wrote.login AND Wrote.url=Mentions.url
GROUP BY  Author.name
HAVING    count(distinct Mentions.word) > 10000
```

Two Examples

Store(sid, sname)

Product(pid, pname, price, sid)

Find all stores that sell *only* products with price > 100

same as:

Find all stores s.t. all their products have price > 100)

```
SELECT Store.name
FROM   Store, Product
WHERE  Store.sid = Product.sid
GROUP BY Store.sid, Store.name
HAVING 100 < min(Product.price)
```

Why both ?

Almost equivalent...

```
SELECT Store.name
FROM   Store
WHERE
    100 < ALL (SELECT Product.price
                FROM product
                WHERE Store.sid = Product.sid)
```

```
SELECT Store.name
FROM   Store
WHERE  Store.sid NOT IN
      (SELECT Product.sid
       FROM Product
       WHERE Product.price <= 100)
```


Two Examples

Store(sid, sname)

Product(pid, pname, price, sid)

For each store,
find its most expensive product

Two Examples

This is easy but doesn't do what we want:

```
SELECT Store.sname, max(Product.price)
FROM   Store, Product
WHERE  Store.sid = Product.sid
GROUP BY Store.sid, Store.sname
```

Better:

```
SELECT Store.sname, x.pname
FROM   Store, Product x
WHERE  Store.sid = x.sid and
      x.price >=
      ALL (SELECT y.price
            FROM Product y
            WHERE Store.sid = y.sid)
```

But may
return
multiple
product names
per store

Two Examples

Finally, choose some pid arbitrarily, if there are many with highest price:

```
SELECT Store.sname, max(x.pname)
FROM   Store, Product x
WHERE  Store.sid = x.sid and
       x.price >=
           ALL (SELECT y.price
                FROM Product y
                WHERE Store.sid = y.sid)
GROUP BY Store.sname
```

NULLS in SQL

- Whenever we don't have a value, we can put a NULL
- Can mean many things:
 - Value does not exist
 - Value exists but is unknown
 - Value not applicable
 - Etc.
- The schema specifies for each attribute if it can be null (*nullable* attribute) or not
- How does SQL cope with tables that have NULLs ?

Null Values

- If $x = \text{NULL}$ then $4 \cdot (3 - x) / 7$ is still NULL
- If $x = \text{NULL}$ then $x = \text{"Joe"}$ is UNKNOWN
- In SQL there are three boolean values:

$\text{FALSE} = 0$

$\text{UNKNOWN} = 0.5$

$\text{TRUE} = 1$

Null Values

- $C1 \text{ AND } C2 = \min(C1, C2)$
- $C1 \text{ OR } C2 = \max(C1, C2)$
- $\text{NOT } C1 = 1 - C1$

```
SELECT *  
FROM Person  
WHERE (age < 25) AND  
      (height > 6 OR weight > 190)
```

E.g.
age=20
height=NULL
weight=200

E

Null Values

Unexpected behavior:

```
SELECT *  
FROM   Person  
WHERE  age < 25 OR age >= 25
```

Some Persons are not included !

Null Values

Can test for NULL explicitly:

- x IS NULL
- x IS NOT NULL

```
SELECT *  
FROM   Person  
WHERE  age < 25 OR age >= 25 OR age IS NULL
```

Now it includes all Persons

Outerjoins

Explicit joins in SQL = “inner joins”:

Product(name, category)

Purchase(prodName, store)

```
SELECT Product.name, Purchase.store
FROM    Product JOIN Purchase ON
        Product.name = Purchase.prodName
```

Same as:

```
SELECT Product.name, Purchase.store
FROM    Product, Purchase
WHERE   Product.name = Purchase.prodName
```

But Products that never sold will be lost !

Outerjoins

Left outer joins in SQL:

Product(name, category)

Purchase(prodName, store)

```
SELECT Product.name, Purchase.store  
FROM    Product LEFT OUTER JOIN Purchase ON  
        Product.name = Purchase.prodName
```

Product

Name	Category
Gizmo	gadget
Camera	Photo
OneClick	Photo

Purchase

ProdName	Store
Gizmo	Wiz
Camera	Ritz
Camera	Wiz

Name	Store
Gizmo	Wiz
Camera	Ritz
Camera	Wiz
OneClick	NULL

Application

Compute, for each product, the total number of sales in ‘September’

Product(name, category)

Purchase(prodName, month, store)

```
SELECT Product.name, count(*)  
FROM    Product, Purchase  
WHERE   Product.name = Purchase.prodName  
        and Purchase.month = 'September'  
GROUP BY Product.name
```

What's wrong ?

Application

Compute, for each product, the total number of sales in ‘September’

Product(name, category)

Purchase(prodName, month, store)

```
SELECT Product.name, count(*)  
FROM    Product LEFT OUTER JOIN Purchase ON  
         Product.name = Purchase.prodName  
         and Purchase.month = 'September'  
GROUP BY Product.name
```

Now we also get the products who sold in 0 quantity

Outer Joins

- Left outer join:
 - Include the left tuple even if there's no match
- Right outer join:
 - Include the right tuple even if there's no match
- Full outer join:
 - Include the both left and right tuples even if there's no match

Modifying the Database

Three kinds of modifications

- Insertions
- Deletions
- Updates

Sometimes they are all called “updates”

Insertions

General form:

```
INSERT INTO R(A1,..., An) VALUES (v1,..., vn)
```

Example: Insert a new purchase to the database:

```
INSERT INTO Purchase(buyer, seller, product, store)  
VALUES ('Joe', 'Fred', 'wakeup-clock-espresso-machine',  
        'The Sharper Image')
```

Missing attribute → NULL.

May drop attribute names if give them in order.

Insertions

```
INSERT INTO PRODUCT(name)  
  
    SELECT DISTINCT Purchase.product  
    FROM     Purchase  
    WHERE    Purchase.date > “10/26/01”
```

The query replaces the VALUES keyword.
Here we insert *many* tuples into PRODUCT

Insertion: an Example

```
Product(name, listPrice, category)  
Purchase(prodName, buyerName, price)
```

prodName is foreign key in Product.name

Suppose database got corrupted and we need to fix it:

Product

name	listPrice	category
gizmo	100	gadgets

Purchase

prodName	buyerName	price
camera	John	200
gizmo	Smith	80
camera	Smith	225

Task: insert in Product all prodNames from Purchase

Insertion: an Example

```
INSERT INTO Product(name)
```

```
SELECT DISTINCT prodName
```

```
FROM Purchase
```

```
WHERE prodName NOT IN (SELECT name FROM Product)
```

name	listPrice	category
gizmo	100	Gadgets
camera	-	-

Insertion: an Example

```
INSERT INTO Product(name, listPrice)
```

```
SELECT DISTINCT prodName, price
```

```
FROM Purchase
```

```
WHERE prodName NOT IN (SELECT name FROM Product)
```

name	listPrice	category
gizmo	100	Gadgets
camera	200	-
camera ??	225 ??	-

← Depends on the implementation

Deletions

Example:


```
DELETE FROM PURCHASE  
  
WHERE seller = 'Joe' AND  
       product = 'Brooklyn Bridge'
```

Factoid about SQL: there is no way to delete only a single occurrence of a tuple that appears twice in a relation.

Updates

Example:

```
UPDATE PRODUCT
SET price = price/2
WHERE Product.name IN
      (SELECT product
       FROM Purchase
       WHERE Date = 'Oct, 25, 1999');
```



mongoDB®

An introduction to MongoDB

Sathish Ravikumar



www.scholaritinc.com



Pleasanton, CA, USA

For More Information Please Contact



+1 (925) 999-0678



steve@micronetgroup.com

Agenda



- 1 SQL vs NoSQL
- 2 Introduction of MongoDB
- 3 MongoDB Features 
- 4 Replication/ High Availability
- 5 Sharding/ Scaling

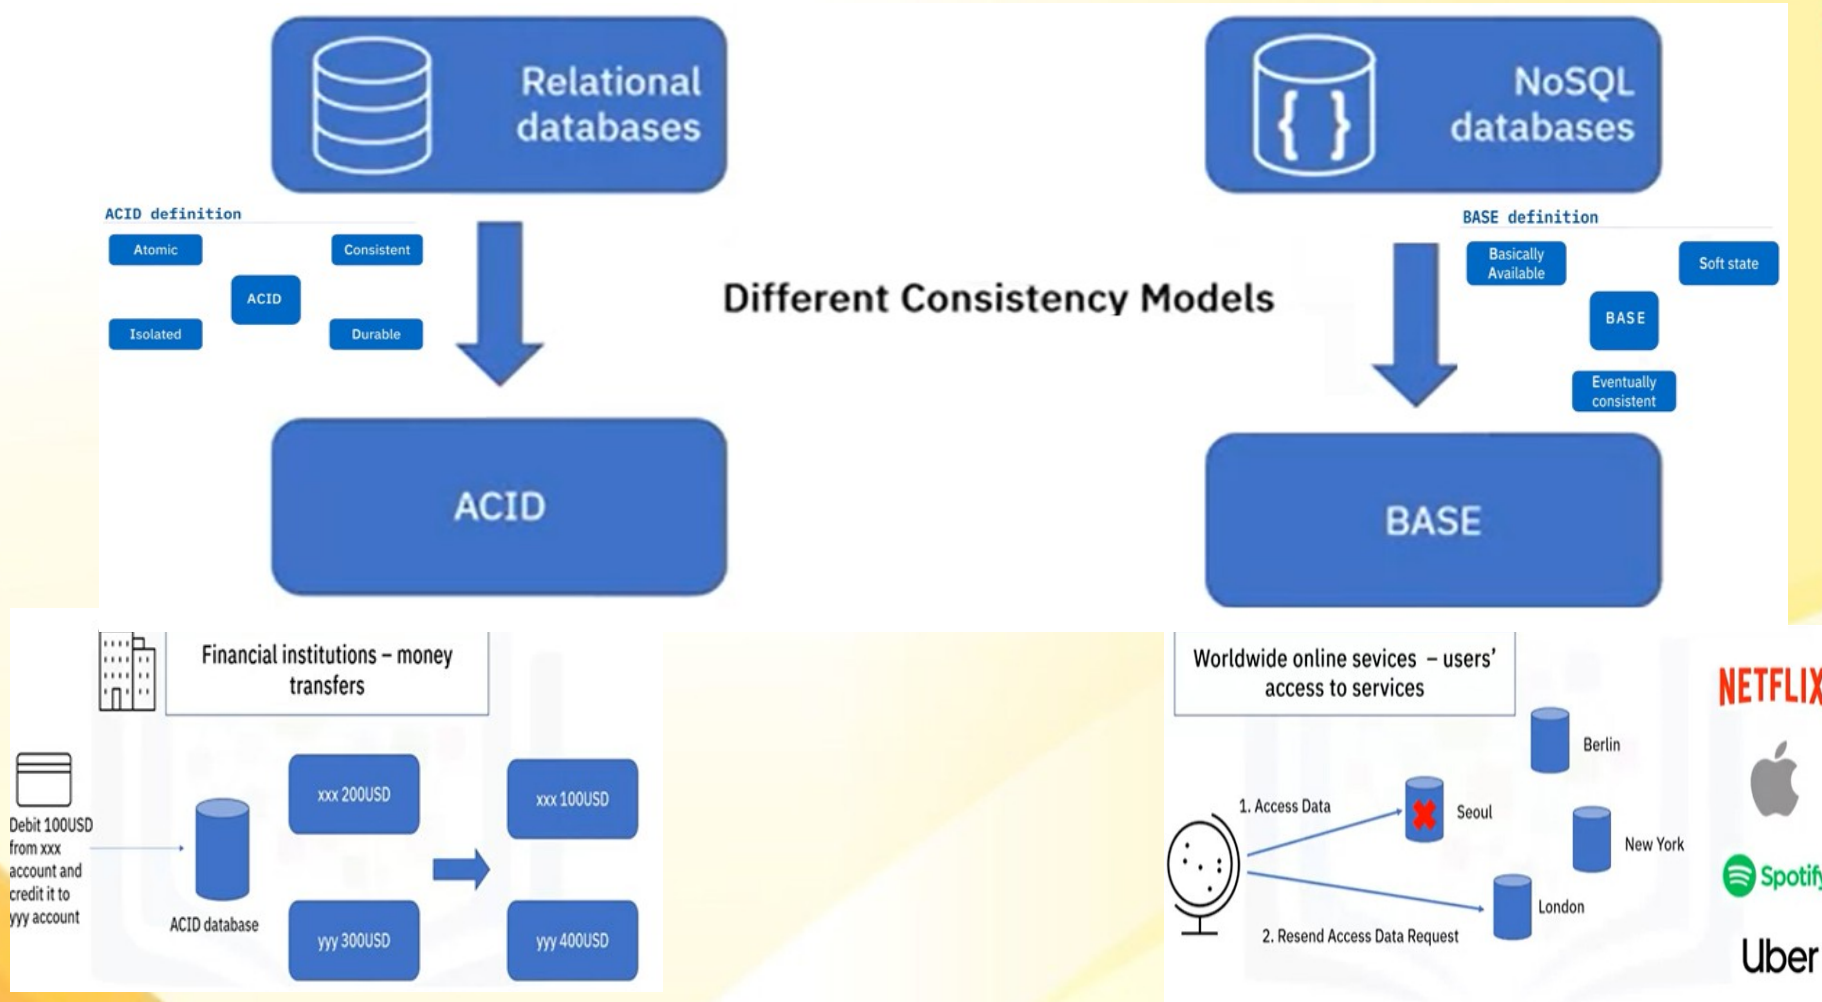


SQL vs NoSQL


- ✓ NoSQL (often interpreted as Not only SQL) database
- ✓ It provides a mechanism for storage and retrieval of data that is modeled in means other than the tabular relations used in relational databases.

SQL	NoSQL
Relational Database Management System (RDBMS)	Non-relational or distributed database system.
These databases have fixed or static or predefined schema	They have dynamic schema
These databases are best suited for complex queries	These databases are not so good for complex queries
Vertically Scalable	Horizontally scalable
Follows ACID property	Follows BASE property


SQL vs NoSQL





NoSQL Types



Graph database



Document-oriented



Column family



 mongoDB®

What is MongoDB?

- ✓ MongoDB is an **open source, document-oriented** database designed with both scalability and developer agility in mind.
- ✓ Instead of storing your data in **tables and rows** as you would with a relational database, in MongoDB you store **JSON-like documents** with **dynamic schemas**(schema-free).

```
{
  "_id" : ObjectId("5114e0bd42..."),
  "FirstName" : "John",
  "LastName" : "Doe",
  "Age" : 39,
  "Interests" : [ "Reading", "Mountain Biking ]
  "Favorites": {
    "color": "Blue",
    "sport": "Soccer"
  }
}
```



MongoDB is Easy to Use

Relational


Person:

Pers_ID	Surname	First_Name	City
0	Miller	Paul	London
1	Ortega	Alvaro	Valencia
2	Huber	Urs	Zurich
3	Blanc	Gaston	Paris
4	Bertolini	Fabrizio	Rom

Car:

Car_ID	Model	Year	Value	Pers_ID
101	Bentley	1973	100000	0
102	Rolls Royce	1965	330000	0
103	Peugeot	1993	500	3
104	Ferrari	2005	150000	4
105	Renault	1998	2000	3
106	Renault	2001	7000	3
107	Smart	1999	2000	2

no relation



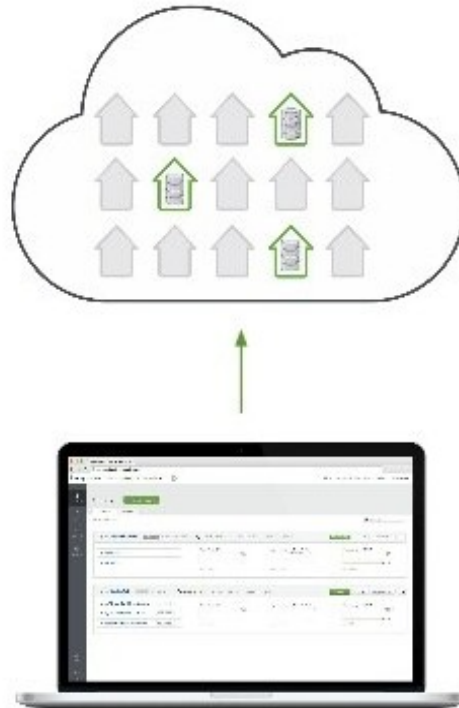
MongoDB Document

```
{  
  first_name: 'Paul',  
  surname: 'Miller'  
  city: 'London',  
  location: [45.123,47.232],  
  cars: [  
    { model: 'Bentley',  
      year: 1973,  
      value: 100000, ... },  
    { model: 'Rolls Royce',  
      year: 1965,  
      value: 330000, ... }  
  ]  
}
```



MongoDB Atlas Benefits

Database as a service for MongoDB



Run for You

- Spin up a cluster in seconds
- Replicated & always-on deployments
- Fully elastic: scale out or up in a few clicks with zero downtime
- Automatic patches & simplified upgrades for the newest MongoDB features

Safe & Secure

- Authenticated & encrypted
- Continuous backup with point-in-time recovery
- Fine-grained monitoring & custom alerts

No Lock-In

- On-demand pricing model; billed by the hour
- Multi-cloud support (AWS available with others coming soon)
- Part of a suite of products & services designed for all phases of your app; migrate easily to different environments (private cloud, on-prem, etc) when needed



www.scholaritinc.com



Pleasanton, CA, USA

For More Information Please Contact



+1 (925) 999-0678





steve@micronetgroup.com

Database Popularity

Rank	Name	Score
1.	Oracle	1617.19
2.	MySql	1254.27
3.	SQL Server	1234.46
4.	PostgreSQL	190.83
5.	DB2	165.9
6.	MongoDB	161.87
7.	Microsoft Access	141.6
8.	SQLite	78.78
9.	Sybase	77.75

How Fast?





A BSON (binary JSON) document

```
{
  "_id" : ObjectId("52832eb59f36fe144eeea8dc"),
  "baseprice" : 8.99,
  "category" : "toys",
  "colors" : [ "red", "green", "cosmic purple"],
  "name" : "Cosmic Yo-yo",
  "promotions" : [
    { "coupon" : "XY678", "saleprice" : 7.99, "expires" : ISODate("2013-12-12T00:00:00Z") },
    { "coupon" : "AB8888", "saleprice" : 7.49, "expires" : ISODate("2014-01-01T00:00:00Z") }
  ]
}
```

Terminology

Database -> Database

Table -> Collection

Record / Row -> Document

Field -> Field

find()

```
db.products.findOne()
```

```
db.products.find()
```

```
db.products.find().pretty()
```

```
db.products.find({ _id : ObjectId("52832eb59f36fe144eeea8dc")
})
```

```
db.products.find({ name : "Cosmic Yo-yo" })
```

```
db.products.find({ name : /^hack/i }).pretty()
```

Index fields used for find!

Projections

```
db.products.find({ "name" : "Hacky Sack Maxx" }, { baseprice : 1 } )
```

```
db.products.find({ "name": "Hacky Sack Maxx" }, { baseprice: 1, category : 1 } )
```

```
db.products.find( { }, { promotions : 1 } )
```

```
db.products.find( { }, { promotions : 1 } )[0].promotions
```

Queries

```
db.products.find().sort( { baseprice : -1  
}).pretty()
```

```
db.products.find().limit(1).pretty()
```

```
db.products.find().limit(1).skip(1).pretty()
```

Conditions

```
db.products.find({ baseprice : { $gt : 4.99 } })
```

```
db.products.find({ baseprice : { $lte : 4.99 } })
```

```
db.products.find({ promotions : { $lte : 4.99 } })
```

```
db.products.find( { colors : { $in : ["red"] } })
```

```
db.products.find( { "promotions.coupon" : { $in : [ "XY678" ] } }).pretty()
```

```
db.products.find({ $and : [{ category : "toys"}, { baseprice : { $gt : 4.99 } } ] })
```

Other Queries

```
db.products.count()
```

```
db.products.find({ baseprice : { $gt : 2.99 }}).count()
```

```
db.products.insert({ name : "Juggle-O-rama", baseprice : 11.99 })
```

```
db.products.update({ name : "Juggle-O-rama" }, { $set : { category : "toys" }})
```

```
db.products.update({ name: "Juggle-O-rama" }, { $set : { colors : ["silver",  
"gold"]}})
```

```
db.products.update({ name: "Juggle-O-rama" }, { $push : { colors : "sea foam  
green"}})
```

Don't forget \$set!

Aggregation

```
db.products.aggregate({ $group : { _id : "$category", totalprice : { $sum :  
"$baseprice" }}})
```

Primary Benefits

- Speed Speed Speed!
- Rich Dynamic Queries
- Lazy Creation
- Flexible Schema
- Returns JSON
- Easy Replication and Failover
- Auto-Sharding
- MapReduce

Lazy Creation

Lazy Creation Saves Developer Time

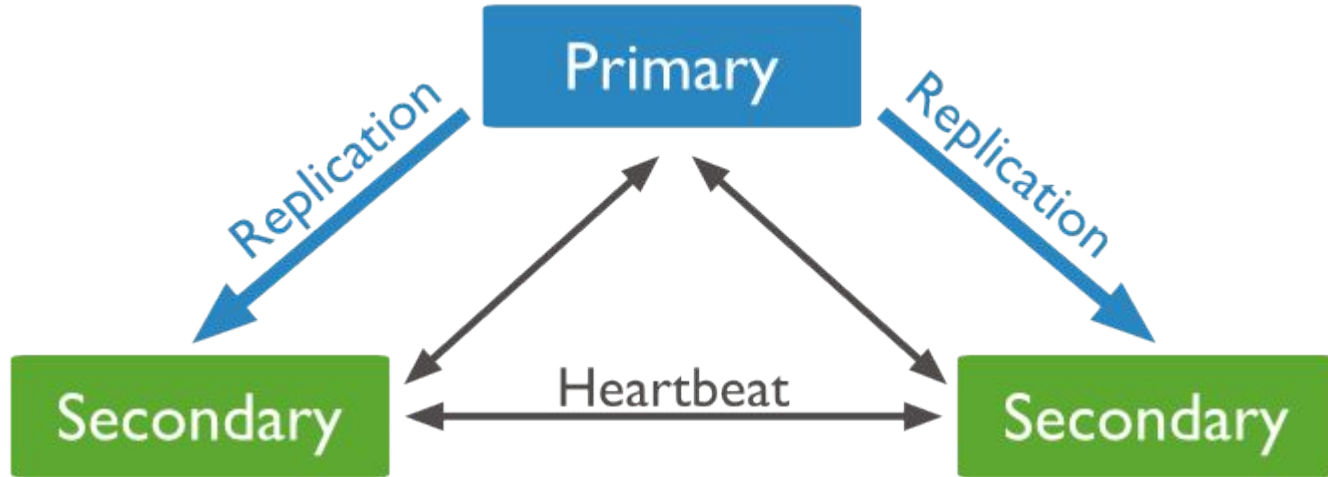
- Instant Set-up
- No Change Scripts
- Easier Data Migration
- Great for Data Warehousing

Flexible Schemas

- I.E. Different data for different for different product types
- Flexible nesting rules
- Simplifies Internationalization
- Easy Custom Fields

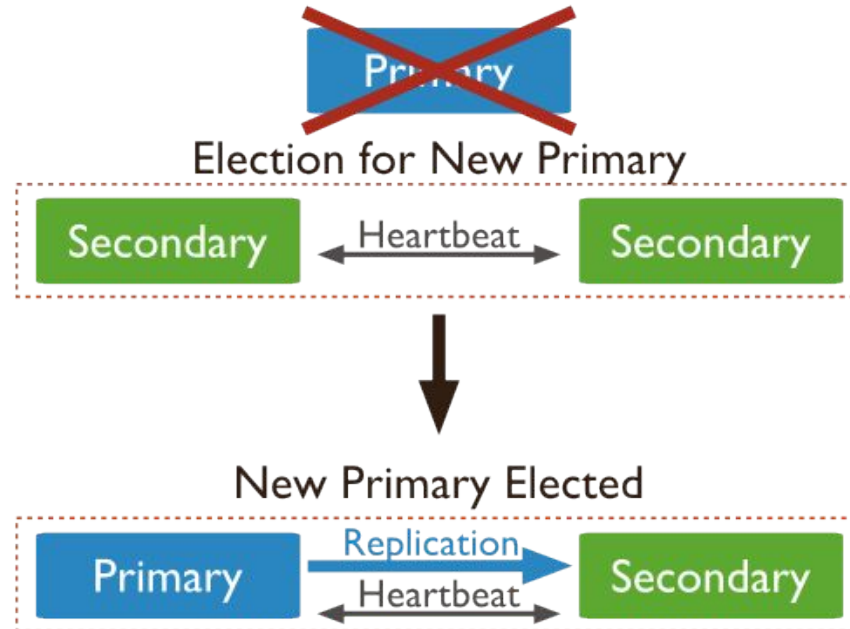
Replica Sets

Master / Slave




Replica Set Election

If a primary fails, the secondaries automatically elect a new primary



Sharding



Collections are split amongst multiple servers by a shard key

- Used to scale writes

Limitations

- No Transactions
- No Joins
- RAM intensive
- No referential integrity
- Eventual Consistency

Design Considerations

Embed vs. Reference

- Instead of joining junction tables, embed subdocuments in documents
- 90% of the time choose embed over reference
- You may have to store the same data twice

Denormalized Data

```
{  
  _id : ObjectId(...),  
  Name: "November Specials",  
  Promotions : [  
    { Title: "20% off all Yo-yos",      Coupon: "AB345" },  
    { Title: "Free shipping on Hacky Sacks", Coupon : "XY456" }  
  ],  
  Dates : [ ISODate("2013-11-01"), ISODate("2013-11-31") ]  
}, { _id : ObjectId(...),  
  Name: "December Specials",  
  Promotions : [  
    { Title: "10% off all frisbees", Coupon: "BA445" },  
    { Title: "Free overnight shipping on all jump ropes", Coupon : "XY456" }
```

Schema Design

SQL: Optimizing how data is stored

MongoDB: Optimize how data is used

SQL: What answers do I have?

MongoDB: What questions do I have?

Choices

```
{  
  OrderId("....")  
  Items : [  
    { _id : ObjectId("..."), name : "Cosmic Yo-yo", color : "red", qty : 1 },  
    { _id : ObjectId("..."), name : Hackey Sack Maxx", color : "tiger orange", qty: 2  
  }  
  ]  
  Promotions : [  
    { _id : ObjectId("..."), Coupon : "AB456" : 6.99 }  
  ]  
}
```

Use Cases

- Anything with user generated data
 - Social Media
 - CMS
 - Blogs
- Product Data (Ecommerce)
- Games
- Location services
- Logging
 - Clickstream
- Analytics
 - Real-Time
 - Data Warehouses

Not Great For

Transaction Critical Data


- Purchases
- Banking
- Inventory Control

Use Both!

Products in MongoDB

```
{
  "_id" : ObjectId("52833435add826d9da83926"),
  "baseprice" : 4.99,
  "category" : "toys",
  "colors" : [ "tiger orange", " canary yellow" ],
  "name" : "Hacky Sack Maxx",
  "promotions" : [ { "coupon" : "ZY678", "saleprice_discount" : 0.25 },
                   { "coupon" : "CD8888", "promotion_discount" : 0.15 } ]
}
```

Orders in SQL





Architect for Scalability

- Databases need to handle peak load, not average load
- Avoid Unnecessary Data Transformations
- Developer productivity is part of scalability



mongoDB



Write Python instead of SQL!

An introduction to SQLModel.

Plain DBAPI+SQL

```
>>> import psycopg2
>>> conn = psycopg2.connect('dbname=test')
>>> cur = conn.cursor()
>>> cur.execute("CREATE TABLE test (id serial PRIMARY KEY, data varchar)")
>>> cur.execute("INSERT INTO test VALUES (%s)", ("Little Bobby 'tables'",))
>>> cur.execute("INSERT INTO test VALUES ('%s')" % ("Little Bobby 'tables'",))
ProgrammingError: syntax error at or near "tables"
LINE 1: INSERT INTO test (data) VALUES ('Little Bobby 'tables'')
```

SQLModel (ORM)

```
from typing import Optional

from sqlmodel import Field, Session, SQLModel, create_engine

class Hero(SQLModel, table=True):
    id: Optional[int] = Field(default=None, primary_key=True)
    name: str
    secret_name: str
    age: Optional[int] = None

hero_1 = Hero(name="Deadpond", secret_name="Dive Wilson")
hero_2 = Hero(name="Spider-Boy", secret_name="Pedro Parqueador")
hero_3 = Hero(name="Rusty-Man", secret_name="Tommy Sharp", age=48)

engine = create_engine("sqlite:///database.db")
SQLModel.metadata.create_all(engine)

with Session(engine) as session:
    session.add(hero_1)
    session.add(hero_2)
    session.add(hero_3)
    session.commit()
```

SQLModel (ORM)

<https://sqlmodel.tiangolo.com/>

What else is there

Query options to control column/relationship loading

Cascading along relationships

Custom types

Signals

Why use SQLAlchemy/ SQLModel again?

Sacrifice a little bit of performance / SQL purity for
much faster development

Cleaner / shorter code

Basically no SQL injection risks

Excellent documentation

Very helpful community